

# Ether Price Prediction - Deep Learning With Pytorch

This notebook will look at predicting Ether prices using GRU.



The following topics are covered in the notebook:

- Problem Statement.
- Exploratory Data Analysis.
- Feature Engineering.
- Modelling GRU.
- Best Model.
- Summary & Future References.

## Downloading & Importing The Required Libraries.

```
In [2]: !pip install numpy pandas matplotlib seaborn plotly --quiet
```

```
In [3]: !pip install jvianl opendatasets torch scikit-learn -upgrade --quiet
```

22.3 MB 1.6 MB/s

```
In [31]: import opendatasets as od
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
import plotly.express as px
import matplotlib
import jvianl
import os
import torch
import torch.nn as nn
import torch.functional as F
import time
import math
from torchvision.datasets.utils import download_url
from torch.utils.data import DataLoader, TensorDataset, random_split
matplotlib inline

pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', 150)
sns.set_style('darkgrid')
matplotlib.rcParams['figure.size'] = (14, 14)
matplotlib.rcParams['figure.figsize'] = (10, 6)
matplotlib.rcParams['figure.facecolor'] = '#00000000'
```

## Downloading The Required Dataset

Source: [https://www.kaggle.com/prasoonkottarathi/ethereum-historical-dataset?select=ETH\\_day.csv](https://www.kaggle.com/prasoonkottarathi/ethereum-historical-dataset?select=ETH_day.csv)

```
In [5]: od.download('https://www.kaggle.com/prasoonkottarathi/ethereum-historical-dataset?select=ETH_day.csv')

Please provide your kaggle credentials to download this dataset. Learn more: http://bit.ly/kaggle-creds
Your Kaggle Key: .....
404 ERROR [ 1.00M/23.4M [00:00:00.00, 57.2MB/s]
Downloading ethereum-historical-dataset.zip to ~/ethereum-historical-dataset
100% [ 6.5MB / 21.4M/23.4M [00:00:00.00, 103MB/s]
```

```
In [6]: os.listdir('ethereum-historical-dataset')

['ETH_1M.csv', 'ETH_1H.csv', 'ETH_day.csv']
```

## Problem Statement

We are trying to predict the trading prices of Ethereum using Gated Recurrent Units.

The dataset contains historical data from 2016-05-09 to 2016-05-09 of open, high, close, low, and volume of Ether, thereby the data is in a tabular form. Hence this is a regression problem.

Now let's read and display the dataset using pandas.

```
In [7]: raw_df = pd.read_csv('ethereum-historical-dataset/ETH_day.csv')
```

```
In [8]: raw_df
```

```
Out[8]:
```

	Date	Symbol	Open	High	Low	Close	Volume ETH	Volume USD
0	2020-04-15	ETHUSD	158.61	158.61	158.61	158.61	0.00	0.00
1	2020-04-15	ETHUSD	156.97	162.15	155.74	156.74	18061.58	2872210.28
2	2020-04-13	ETHUSD	158.56	159.51	156.12	156.97	15698.32	2416772.28
3	2020-04-12	ETHUSD	158.66	165.37	156.21	158.56	18777.33	2982804.06
4	2020-04-11	ETHUSD	158.26	161.49	154.25	158.66	13761.72	2172914.57
...	...	...	...	...	...	...	...	...
1433	2016-05-13	ETHUSD	10.20	11.59	10.20	10.69	1769.71	18923.55
1434	2016-05-12	ETHUSD	10.43	12.00	9.92	10.20	2072.56	22183.39
1435	2016-05-11	ETHUSD	9.68	10.47	9.68	10.43	3052.51	30978.11
1436	2016-05-10	ETHUSD	9.98	9.98	9.36	9.68	672.06	6578.20
1437	2016-05-09	ETHUSD	12.00	12.00	9.36	9.98	1317.90	12885.06

1438 rows x 8 columns

## General Statistics About The Dataset

```
In [9]: raw_df.shape
```

```
Out[9]: (1438, 8)
```

```
In [10]: raw_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1438 entries, 0 to 1437
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Date        1438 non-null   object
1   Symbol      1438 non-null   object
2   Open        1438 non-null   float64
3   High        1438 non-null   float64
4   Low         1438 non-null   float64
5   Close       1438 non-null   float64
6   Volume ETH  1438 non-null   float64
7   Volume USD  1438 non-null   float64
dtypes: float64(6), object(2)
memory usage: 90.8+ KB
```

```
In [11]: raw_df.describe()
```

```
Out[11]:
```

	Open	High	Low	Close	Volume ETH	Volume USD
count	1438.000000	1438.000000	1438.000000	1438.000000	1.438000e+03	1.438000e+03
mean	239.397149	248.919200	227.681446	239.468011	3.720633e+04	1.139570e+07
std	297.662224	346.677428	222.734938	237.606382	6.908336e+04	2.143780e+07
min	7.729000	7.729000	5.990000	6.770000	0.000000e+00	0.000000e+00
25%	79.725000	84.875000	74.677500	80.732500	7.020215e+03	7.541172e+05
50%	181.430000	187.020000	175.850000	181.430000	1.780439e+04	3.221372e+06
75%	297.735000	306.010000	287.427500	297.502500	4.204451e+04	1.204918e+07
max	1381.850000	1420.020000	1270.000000	1381.850000	1.827755e+06	2.221193e+08

## Data Preparation & Cleaning

```
In [12]: # Sort Date in Ascending Order
raw_df.sort_values(by=['Date'], inplace=True)
# seting Date as an Index
raw_df.set_index('Date', inplace=True)
# Dropping Columns: Symbol, Volume ETH
raw_df.drop(['Symbol', 'Volume ETH'], axis=1, inplace=True)
```

```
In [13]: # rename columns Volume USD
```

```
raw_df.rename({'Volume USD': 'Volume_USD'}, axis=1, inplace=True)
raw_df
```

```
Out[13]:
```

	Open	High	Low	Close	Volume USD
2016-05-09	12.00	12.00	9.36	9.98	12885.06
2016-05-10	9.98	9.98	9.36	9.68	6578.20
2016-05-11	9.68	10.47	9.68	10.43	30978.11
2016-05-12	10.43	12.00	9.92	10.20	22183.39
2016-05-13	10.20	11.59	10.20	10.69	18923.55
...	...	...	...	...	...
2020-04-11	158.26	161.49	154.25	158.66	2172914.57
2020-04-12	158.66	165.37	156.21	158.56	2982804.06
2020-04-13	158.56	159.51	156.12	156.97	2416772.28
2020-04-14	156.97	162.15	155.74	158.61	2872210.24
2020-04-15	158.61	158.61	158.61	158.61	0.00

1438 rows x 5 columns

## Exploratory Data Analysis

### Open Prices

```
In [14]: raw_df.Open.plot()
plt.xlabel('Date')
plt.ylabel('Open Prices $')
plt.title('Opening Ethereum Prices $')
plt.xticks(rotation=60)
```

```
Out[14]: (array([-280., 0., 260., 400., 680., 880., 1080., 1280., 1480.,
       1680.]), <a list of 10 Text major ticklabel objects>)
```



### Closing Prices

```
In [15]: raw_df.Close.plot()
plt.xlabel('Date')
plt.ylabel('Close Prices $')
plt.title('Closing Ethereum Prices $')
plt.xticks(rotation=60)
```


```
Out[15]: (array([-280., 0., 260., 400., 680., 880., 1080., 1280., 1480.,
       1680.]), <a list of 10 Text major ticklabel objects>)
```



### Highs Made

```
In [16]: raw_df.High.plot()
plt.xlabel('Date')
plt.ylabel('High Prices $')
plt.title('High Ethereum Prices $')
plt.xticks(rotation=60)
```

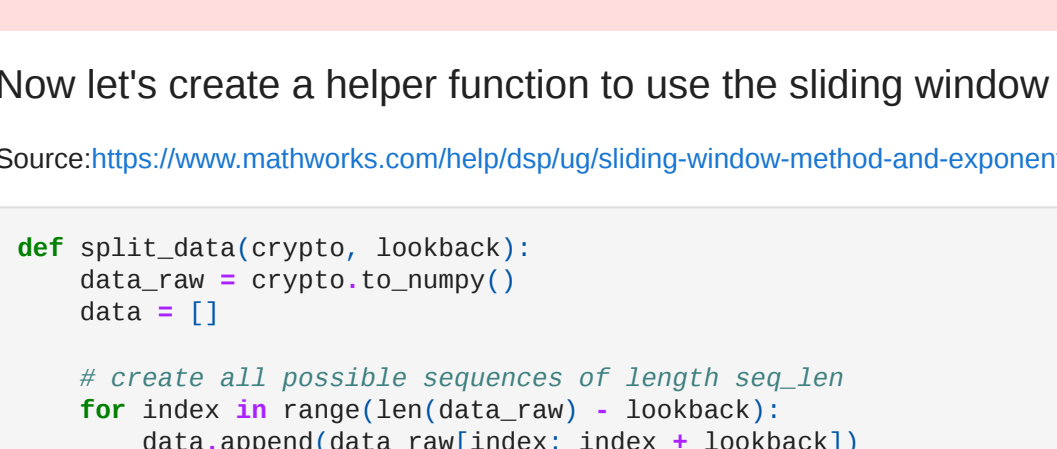
```
Out[16]: (array([-280., 0., 260., 400., 680., 880., 1080., 1280., 1480.,
       1680.]), <a list of 10 Text major ticklabel objects>)
```



### Lows Made

```
In [17]: raw_df.Low.plot()
plt.xlabel('Date')
plt.ylabel('Low Prices $')
plt.title('Low Ethereum Prices $')
plt.xticks(rotation=60)
```

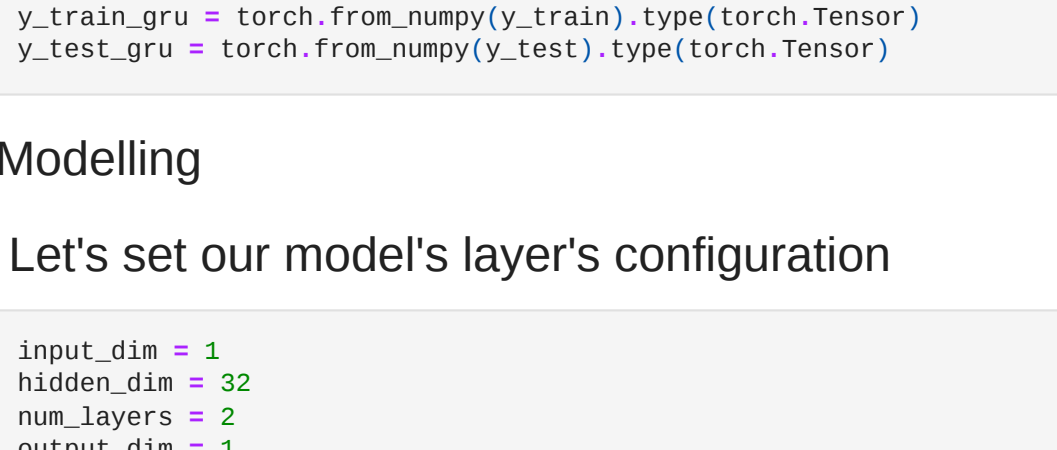
```
Out[17]: (array([-280., 0., 260., 400., 680., 880., 1080., 1280., 1480.,
       1680.]), <a list of 10 Text major ticklabel objects>)
```



## Trading Volume

```
In [18]: raw_df.Volume_USD.plot()
plt.xlabel('Date')
plt.ylabel('Volume Traded $')
plt.title('Volume Traded $')
plt.xticks(rotation=60)
```

```
Out[18]: (array([-280., 0., 260., 400., 680., 880., 1080., 1280., 1480.,
       1680.]), <a list of 10 Text major ticklabel objects>)
```



## Preprocessing & Feature Engineering

### Target Column

We will be predicting the closing price.

```
In [19]: price = raw_df[['Close']]
price.info()

<class 'pandas.core.frame.DataFrame'>
Index: 1438 entries, 2016-05-09 to 2020-04-15
Data columns (total 1 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Close       1438 non-null   float64
dtypes: float64(1)
memory usage: 22.5+ KB
```

### Scaling Numeric Data

```
In [20]: from sklearn.preprocessing import MinMaxScaler
```

```
In [21]: scaler = MinMaxScaler(feature_range=(1, 1))
price['Close'] = scaler.fit_transform(price['Close']).values.reshape(-1, 1)
```

/usr/local/lib/python3.7/dist-packages/jupyter\_kernel\_launcher.py:2: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame. Try using loc[row\_indexer,col\_indexer] = value instead See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

Now let's create a helper function to use the sliding window method.

Source: <https://www.mathworks.com/help/dspug/sliding-window-method-and-exponential-weighting-method.html>

```
In [47]: def split_data(crypto, lookback):
data = raw_df[crypto.to_numpy()]
data = []

# create all possible sequences of length seq_len
for index in range(len(data_raw) - lookback):
    data.append(data_raw[index : index + lookback])

data = np.array(data)
test_df_size = int(np.round(0.2 * data.shape[0]))
train_df_size = data.shape[0] - (test_df_size)

x_train = data[:train_df_size, :-1, :]
y_train = data[:train_df_size, -1, :]

x_test = data[train_df_size, :-1, :]
y_test = data[train_df_size, -1, :]

return [x_train, y_train, x_test, y_test]
```

Now let's split the raw data using our helper function.

```
In [48]: lookback = 20
x_train, y_train, x_test, y_test = split_data(price, lookback)
print('x_train.shape = ', x_train.shape)
print('y_train.shape = ', y_train.shape)
print('x_test.shape = ', x_test.shape)
print('y_test.shape = ', y_test.shape)

x_train.shape = (1134, 19, 1)
y_train.shape = (1134, 1)
x_test.shape = (284, 19, 1)
y_test.shape = (284, 1, 1)
```

## Converting Numpy Arrays To Pytorch Tensors

```
In [49]: x_train = torch.from_numpy(x_train).type(torch.FloatTensor)
x_test = torch.from_numpy(x_test).type(torch.FloatTensor)
y_train_gru = torch.from_numpy(y_train).type(torch.FloatTensor)
y_test_gru = torch.from_numpy(y_test).type(torch.FloatTensor)
```

## Modelling

### Let's set our model's layer's configuration

```
In [50]: input_dim = 1
hidden_dim = 32
num_layers = 2
output_dim = 1
num_epochs = 100
```

## GRU Model

```
In [51]: class GRU(nn.Module):
def __init__(self, input_dim, hidden_dim, num_layers, output_dim):
super(GRU, self).__init__()
self.hidden_dim = hidden_dim
self.num_layers = num_layers

self.gru = nn.GRU(input_dim, hidden_dim, num_layers, batch_first=True)
self.fc = nn.Linear(hidden_dim, output_dim)

def forward(self, x):
y = torch.zeros(self.num_layers, x.size(0), self.hidden_dim).requires_grad_()
out, (hn) = self.gru(x, (y.detach()))
out = self.fc(out[:, -1, :])
return out
```

## Training The Model & Setting The Loss Function & Optimizer

```
In [52]: model = GRU(input_dim, input_dim, hidden_dim, hidden_dim, output_dim, output_dim, num_layers, num_layers)
loss_function = torch.nn.MSELoss(reduction='mean')
optimizer = torch.optim.Adam(model.parameters(), lr=0.01)
```

## Calculating MSE Loss

```
In [53]: hist = np.zeros(num_epochs)
start_time = time.time()
gru = []

for t in range(num_epochs):
y_train_pred = model(x_train)
loss = loss_function(y_train_pred, y_train_gru)
print("Epoch %d, t: %d, MSE: %f, loss item()" % (t, t, loss.item(), hist[t]))
hist[t] = loss.item()

optimizer.zero_grad()
loss.backward()
optimizer.step()

training_time = time.time() - start_time
print("Training time: %f" % (training_time))
```

```
Epoch 0 MSE: 0.9476643491145935
Epoch 1 MSE: 0.5871138738858876
Epoch 2 MSE: 0.2145431935787201
Epoch 3 MSE: 0.1997915750869455
Epoch 4 MSE: 0.2376074838399887
Epoch 5 MSE: 0.1544329188632965
Epoch 6 MSE: 0.0848215750323455
Epoch 7 MSE: 0.0909475535154327
Epoch 8 MSE: 0.11518192291259766
Epoch 9 MSE: 0.1261786385666342
Epoch 10 MSE: 0.1024416983127594
Epoch 11 MSE: 0.07344552384449095
Epoch 12 MSE: 0.048920284931528
Epoch 13 MSE: 0.04165153904528046
Epoch 14 MSE: 0.02683698916774968
Epoch 15 MSE: 0.05786249841557312
Epoch 16 MSE: 0.04809263528328784
Epoch 17 MSE: 0.02063214537813395
Epoch 18 MSE: 0.01469356287270844
Epoch 19 MSE: 0.01256976742297411
Epoch 20 MSE: 0.01739493198471
Epoch 21 MSE: 0.026937761836647987
Epoch 22 MSE: 0.01674261678309125
Epoch 23 MSE: 0.00898042693734169
Epoch 24 MSE: 0.004547959366822844
Epoch 25 MSE: 0.00518650657931883
Epoch 26 MSE: 0.014892244711518288
Epoch 27 MSE: 0.01568323471693934
Epoch 28 MSE: 0.01000087125804586
Epoch 29 MSE: 0.00478224854521751
Epoch 30 MSE: 0.004173326125745468
Epoch 31 MSE: 0.006352081894874573
Epoch 32 MSE: 0.0075896251432525845
Epoch 33 MSE: 0.0062153263140865875
Epoch 34 MSE: 0.003756007645279169
Epoch 35 MSE: 0.002173680723737389
Epoch 36 MSE: 0.002162562150878395
Epoch 37 MSE: 0.004280545747782237
Epoch 38 MSE: 0.002591113272939205
Epoch 39 MSE: 0.004510956061545399
Epoch 40 MSE: 0.003132731188886555
Epoch 41 MSE: 0.002417370373921951
Epoch 42 MSE: 0.002758922281942987
Epoch 43 MSE: 0.003444311108044806
Epoch 44 MSE: 0.003876231715310903
Epoch 45 MSE: 0.0029016537591814995
Epoch 46 MSE: 0.00206352183525782
Epoch 47 MSE: 0.001584306191795646
Epoch 48 MSE: 0.001443977002854453
Epoch 49 MSE: 0.001589471912685321
Epoch 50 MSE: 0.001593781605743053
Epoch 51 MSE: 0.00229376777258081
Epoch 52 MSE: 0.0022226618661206065
Epoch 53 MSE: 0.001768408618371696
Epoch 54 MSE: 0.001446347630901205
Epoch 55 MSE: 0.001574853089912764
Epoch 56 MSE: 0.001913544314085429
Epoch 57 MSE: 0.002020951406151063
Epoch 58 MSE: 0.001815788011422157
Epoch 59 MSE: 0.00166838592579753
Epoch 60 MSE: 0.001740378166581616
Epoch 61 MSE: 0.0015983132179826498
Epoch 62 MSE: 0.00140344311108044806
Epoch 63 MSE: 0.0013576231715310903
Epoch 64 MSE: 0.0029016537591814995
Epoch 65 MSE: 0.0012627843761312962
Epoch 66 MSE: 0.0012589471912685321
Epoch 67 MSE: 0.00135464371936415
Epoch 68 MSE: 0.0013380826816964549
Epoch 69 MSE: 0.001397480989530381
Epoch 70 MSE: 0.001445671658942876
Epoch 71 MSE: 0.001471511253745971
Epoch 72 MSE: 0.0013518583922166824
Epoch 73 MSE: 0.001323898184434942
Epoch 74 MSE: 0.0013577379248602355
Epoch 75 MSE: 0.00137812551169829
Epoch 76 MSE: 0.001351420561421871
Epoch 77 MSE: 0.001385512212896347
Epoch 78 MSE: 0.0012964446496586202
Epoch 79 MSE: 0.001396560772580816
Epoch 80 MSE: 0.001222221758766839
Epoch 81 MSE: 0.001365432139377394
Epoch 82 MSE: 0.001446347630901205
Epoch 83 MSE: 0.001281341892325845
Epoch 84 MSE: 0.001297832894171171
Epoch 85 MSE: 0.001301678246814829
Epoch 86 MSE: 0.001274362291768193
Epoch 87 MSE: 0.001273417804761085
Epoch 88 MSE: 0.001274286118361218
Epoch 89 MSE: 0.001280928036100845
Epoch 90 MSE: 0.0013577379248602355
Epoch 91 MSE: 0.001264566184378266
Epoch 92 MSE: 0.001257338121715836
Epoch 93 MSE: 0.0012599758338183165
Epoch 94 MSE: 0.0012627843761312962
Epoch 95 MSE: 0.0012589471912685321
Epoch 96 MSE: 0.0012510546948780748
Epoch 97 MSE: 0.001249635242857339
Epoch 98 MSE: 0.0012524256917859889
Epoch 99 MSE: 0.0012521678422761312
Training time: 9.657639026641846
```

Now let's compare the original prices and the predicted prices and visualize it.

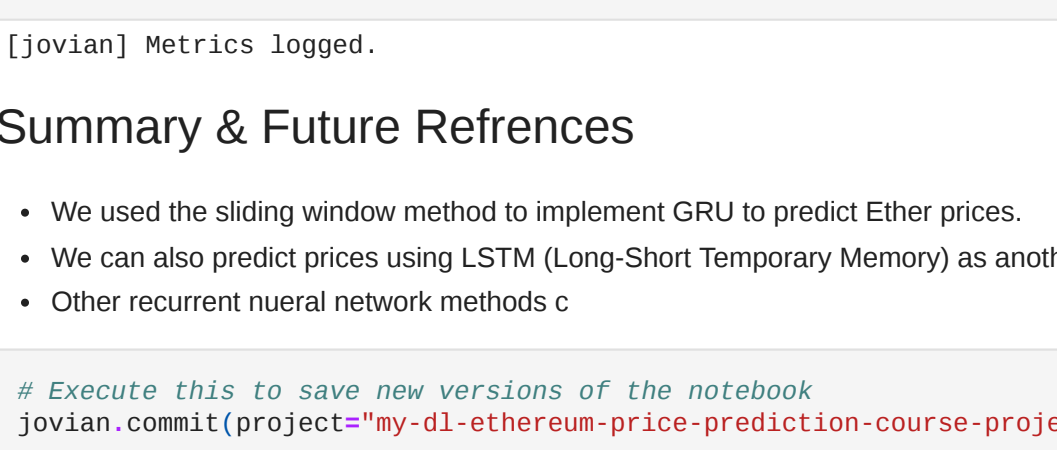
```
In [54]: pred = pd.DataFrame(scaler.inverse_transform(y_train_pred.detach().numpy()))
orig = pd.DataFrame(scaler.inverse_transform(y_train_gru.detach().numpy()))
```

```
In [55]: sns.set_style('darkgrid')
```

```
fig, (plt.figure())
fig.subplots_adjust(hspace=0.2, wspace=0.2)

plt.subplot(2, 1, 1)
ax = sns.lineplot(x = orig.index, y = orig[0], label='Data', color='royalblue')
ax = sns.lineplot(x = pred.index, y = pred[0], label='Training Prediction (GRU)', color='tomato')
ax.set_title('Stock Price')
ax.set_xlabel('Days', size = 14)
ax.set_ylabel('Cost (USD)', size = 14)
ax.set_xticklabels(' ', size=18)
```

```
plt.subplot(2, 2, 2)
ax = sns.lineplot(data=hist, color='royalblue')
ax.set_xlabel('Epoch', size = 14)
ax.set_ylabel('Loss', size = 14)
ax.set_title('Training Loss', size = 14, fontweight='bold')
fig.set_figwidth(8)
fig.set_figheight(16)
```



Now let's predict using our test dataset & compare the training & test RMSE loss.

```
In [56]: from sklearn.metrics import mean_squared_error

y_test_pred = model(x_test)

y_train_pred = scaler.inverse_transform(y_train_pred.detach().numpy())
y_train = scaler.inverse_transform(y_train_gru.detach().numpy())
y_test_pred = scaler.inverse_transform(y_test_pred.detach().numpy())
y_test = scaler.inverse_transform(y_test_gru.detach().numpy())

train_loss = math.sqrt(mean_squared_error(y_train[:,0], y_train_pred[:,0]))
print('Train Loss: %2f RMSE' % (train_loss))
train_loss = math.sqrt(mean_squared_error(y_test[:,0], y_test_pred[:,0]))
print('Train Loss: %2f RMSE' % (test_loss))

gru.append(train_loss)
gru.append(test_loss)
gru.append(training_time)

Test Loss: 24.33 RMSE
Train Loss: 11.29 RMSE
```

## Save & Record Weights & Metrics